

r_tutorial_week_02

Werkcollege 2

Sebastian A. J. Lortz

23 June 2026

Table of contents

How to use this tutorial	2
1 Start workflow	2
Your turn	3
2 Getting the data ready	3
Your turn	3
3 Loading the data	4
Your turn	4
4 Data frames	4
Your turn	5
5 Looking at your data	6
Your turn	7
6 Cleaning and changing your data	8
6.1 Filtering: choosing rows	8
Your turn	8
6.2 Choosing and removing columns	8
Your turn	8
6.3 Making a new variable	9
Your turn	9
6.4 Recoding values	9
Your turn	10
6.5 Putting values into categories	10
Your turn	10

6.6	Taking a random subsample	10
	Your turn	11
7	Describing your sample	11
7.1	Counting categories with table()	11
7.2	Numbers in short: mean and sd	11
7.3	Histograms with hist()	12
	Your turn	12
7.4	Boxplots with boxplot()	12
	Your turn	14
8	Saving your cleaned data	14
	Your turn	15
9	Recap: what you can now do	15

How to use this tutorial

Welcome back! Last week you learned some basics including objects, vectors, indexing, and functions. This week we work with a **dataset**: a whole table of data, like the answers from our class questionnaire.

Work through this page carefully, from top to bottom. Every short explanation is followed by a **Your turn** task to do yourself.

General tips:

- Read each part slowly and in order. Each step builds on the one before it.
- Type the code yourself instead of copying and pasting it. You learn faster this way.
- Errors are normal. Read the message and try again.
- If you are stuck, ask a teacher.
- Save your work often.

1 Start workflow

Based on last week's tutorial, we **always** start with the following **start workflow**:

1. Open **RStudio**: click on RStudio application.
2. Set the **working directory**: click *Session > Set Working Directory > Choose Directory*.

3. Open a **new script**: click *File > New File > R Script*. Or, open an **existing script**: click *File > Open File >* and navigate to the file.
4. On the first line, write a comment with the title, for example: `# R tutorial - Week 2`.
5. Save the script as `week2.R` in your `research_methods` folder.

Today we work in a **new script**. However, often we need to pick-up previous work and open an existing script.

Your turn

Complete the **start workflow** for working with a new script.

2 Getting the data ready

Our class answered a short questionnaire. The answers were collected, made **anonymous** (no names, only a number per person), and saved as a **CSV file**. A CSV file is a simple table where the values are separated by commas. Almost every program can open it.

Your turn

1. Download the file `class_survey.csv` from **Brightspace**.
2. Move it into the `data` folder inside your `research_methods` folder (the same `data` folder from the structure we made in Week 1).

Why a data folder?

Keeping your data in one `data` folder, separate from your scripts, keeps your project tidy and makes the file easy to find later.

3 Loading the data

RStudio does not know about the file until you load it. We load a CSV with the function `read.csv()`. Inside the brackets we put the **path** to the file, in quotes. Because our working directory is `research_methods`, and the file is in the `data` folder, the path is `"data/class_survey.csv"`.

We store the table in an object so we can use it again. We will call it `survey`.

```
survey <- read.csv("data/class_survey.csv")
```

The table is now stored. Look at the **Environment** pane (top-right): `survey` appears, with the number of rows (people) and columns (questions).

Watch out

If you see `Error: cannot open file ... No such file or directory`, R cannot find the file. Check two things:

1. your working directory is `research_methods`: run `getwd()` to print the current working directory
2. the file really sits in the `data` folder and is spelled exactly, including `.csv`: in the bottom-right pane you can click on files and navigate to the folder.

Going further (optional): Excel files

Other file formats, such as Excel files (`.xlsx` or `.xls`), cannot be opened with `read.csv()`. They need an extra **package**, for example `readxl` with its function `read_excel()`. The easiest path for now is to open the Excel file once and save it as a CSV.

Your turn

Load the data into an object called `survey`, then check that `survey` appears in the Environment pane.

4 Data frames

The object `survey` is a **data frame**. A data frame is R's word for a table, in our case:

- each **row** is one participant (one person who answered);
- each **column** is one **variable** (one question or measure).

You can think of a data frame as several vectors (from Week 1) standing side by side, one per column.

To use a single column, write the data frame name, then a **dollar sign** \$, then the column name:

```
survey$age    # all the ages, as a vector
#>  [1] 25 17 22 20 21 18 24 19 17 22 18 18 21 24 21 19 19 19 20 24 18 18 25 21 25
#> [26] 18 25 23 18 21
```

A column is just a vector, so everything you learned last week still works:

```
mean(survey$age)      # average age
#>  [1] 20.66667
survey$age[1]         # the first person's age
#>  [1] 25
survey$age[survey$age > 20] # ages above 20
#>  [1] 25 22 21 24 22 21 24 21 24 25 21 25 25 23 21
```

You can also pick a value by its row and column, inside square brackets, with a **comma** in between: `survey[row, column]`.

```
survey[1, ]    # the whole first row (a comma, then nothing = all columns)
#>   id age self_efficacy study_hours program
#> 1  1  25           40           8         1
survey[1, 2]   # row 1, column 2
#>  [1] 25
```

Your turn

Print the `self_efficacy` column. Then print the `self_efficacy` score of the third person, using \$ and square brackets. Also, print the minimum and maximum `self_efficacy` score.

5 Looking at your data

Before changing anything, always look at your data first. Here are the most useful commands.

How big is the table?

```
dim(survey)      # dimensions: number of rows and columns
#> [1] 30  5
nrow(survey)     # number of rows (people)
#> [1] 30
ncol(survey)     # number of columns (variables)
#> [1] 5
```

What are the columns called?

```
names(survey)
#> [1] "id"           "age"           "self_efficacy" "study_hours"
#> [5] "program"
```

See the first few rows:

```
head(survey)
#>   id age self_efficacy study_hours program
#> 1  1  25           40           8       1
#> 2  2  17           31          18       2
#> 3  3  22           24           6       1
#> 4  4  20           31          14       3
#> 5  5  21           35          13       1
#> 6  6  18           30          16       2
```

See the **type** of each variable with **str()** (short for *structure*):

```
str(survey)
#> 'data.frame':   30 obs. of  5 variables:
#>  $ id          : int  1 2 3 4 5 6 7 8 9 10 ...
#>  $ age         : int  25 17 22 20 21 18 24 19 17 22 ...
#>  $ self_efficacy: int  40 31 24 31 35 30 38 34 33 40 ...
#>  $ study_hours  : int   8 18 6 14 13 16 7 2 14 12 ...
#>  $ program      : int   1 2 1 3 1 2 2 3 3 1 ...
```

Notice that `program` is shown as a number (`int`). R does not know yet that it is really a **category** (1, 2, or 3 stand for study programmes). We will fix that soon.

Get a quick summary of every variable with `summary()`:

```
summary(survey)
#>      id      age      self_efficacy      study_hours
#>  Min.   : 1.00   Min.   :17.00   Min.   :22.00   Min.   : 2.00
#> 1st Qu.: 8.25   1st Qu.:18.00   1st Qu.:29.25   1st Qu.: 8.00
#>  Median :15.50   Median :20.50   Median :33.00   Median :13.50
#>  Mean   :15.50   Mean   :20.67   Mean   :33.23   Mean   :12.17
#> 3rd Qu.:22.75   3rd Qu.:22.75   3rd Qu.:38.00   3rd Qu.:16.00
#>  Max.   :30.00   Max.   :25.00   Max.   :44.00   Max.   :22.00
#>      program
#>  Min.   :1.000
#> 1st Qu.:1.000
#>  Median :2.000
#>  Mean   :2.033
#> 3rd Qu.:3.000
#>  Max.   :3.000
```

For numbers, `summary()` shows the smallest value, the 1st Quartile, the median, the average, the 3rd Quartile, and the largest value. Because R still thinks `program` is a number, it even shows its “average”, which makes no sense for a category. Another reason to recode it later.

Finally, you can open the table in a spreadsheet view. Type this in RStudio (the capital **V** matters):

```
View(survey)
```

Watch out

It is `View()` with a capital **V**. Remember from Week 1: R is case-sensitive, so `view()` with a small `v` will not work.

Your turn

Use three of these commands on `survey`: `dim()`, `names()`, and `summary()`. How many people answered the questionnaire? How many variables are in the dataset? What’s the average age?

6 Cleaning and changing your data

Real data almost always needs some tidying before we describe it. We will keep the loaded data in `survey` and change it step by step.

6.1 Filtering: choosing rows

Filtering means keeping only the rows that match a rule. We use square brackets with a condition before the comma: `survey[condition, ]`. The empty part after the comma means “keep all columns”.

```
adults <- survey[survey$age >= 18, ] # keep only people aged 18 or older
nrow(adults)                         # how many are left?
#> [1] 28
```

Read it as: “keep the rows of `survey` where `age` is 18 or higher.”

Watch out

Do not forget the **comma**. `survey[survey$age >= 18, ]` keeps rows. Without the comma, R gets confused, because it cannot tell rows from columns.

Your turn

Create an object called `keen` that keeps only the people who study more than 15 hours per week (`study_hours > 15`). How many people are in `keen`?

6.2 Choosing and removing columns

Sometimes a column is not useful for analysis. The `id` number, for example, only tells people apart; it is not something we study. We remove a column by setting it to `NULL`:

```
survey$id <- NULL # remove the id column
names(survey)     # id is gone
#> [1] "age"          "self_efficacy" "study_hours"   "program"
```

Your turn

Look at the column names again with `names(survey)` and check that `id` is no longer there. What other function could you use to check this?

6.3 Making a new variable

We can build a new column from the columns we already have. Give the new column a name with \$, and put a calculation on the right of the arrow <-.

```
survey$study_per_day <- survey$study_hours / 7 # weekly hours -> daily hours
head(survey)
#>   age self_efficacy study_hours program study_per_day
#> 1  25             40           8      1      1.1428571
#> 2  17             31          18      2      2.5714286
#> 3  22             24           6      1      0.8571429
#> 4  20             31          14      3      2.0000000
#> 5  21             35          13      1      1.8571429
#> 6  18             30          16      2      2.2857143
```

The new column `study_per_day` now appears in the table.

Your turn

Make a new variable `study_per_day_rounded` that rounds `study_per_day` to whole numbers, using the function `round()`. Then look at the first rows.

6.4 Recoding values

Recoding means changing values into clearer ones. Our `program` column holds the codes 1, 2, and 3, but those numbers are hard to read. We will make a new column with real names. We start it as NA (R's word for "not available / empty") and then fill it in:

```
survey$program_name <- NA # a new, empty column
survey$program_name[survey$program == 1] <- "Pedagogy"
head(survey)
#>   age self_efficacy study_hours program study_per_day program_name
#> 1  25             40           8      1      1.1428571    Pedagogy
#> 2  17             31          18      2      2.5714286         <NA>
#> 3  22             24           6      1      0.8571429    Pedagogy
#> 4  20             31          14      3      2.0000000         <NA>
#> 5  21             35          13      1      1.8571429    Pedagogy
#> 6  18             30          16      2      2.2857143         <NA>
```

Read one of those lines as: “for the rows where `program` equals 1, set `program_name` to Pedagogy.” We kept the original `program` column, which is good practice: never throw away the raw values.

 Watch out

To test “is equal to”, use **two** equals signs `==`, just like in Week 1. One equals sign `=` will not work here.

Your turn

Check your recoding with `table(survey$program_name)`. You should see the pedagogy programme names and NAs with a count for each. Additionally, recode psychology (value 2) and sociology (value 3) into the `program_name` variable.

6.5 Putting values into categories

Categorizing means turning a number into groups. A common tool is `ifelse()`. It takes three parts: a test, the value if the test is `TRUE`, and the value if it is `FALSE`.

```
survey$efficacy_level <- ifelse(survey$self_efficacy >= 32, "higher", "lower")
table(survey$efficacy_level)
#>
#> higher lower
#>      16    14
```

Here, everyone with a self-efficacy score of 32 or more gets “higher”, and the rest get “lower”.

Your turn

Make a variable `studies_a_lot` that is “yes” when `study_hours` is 14 or more, and “no” otherwise. Then count the groups with `table()`.

6.6 Taking a random subsample

A **subsample** is a smaller, random part of your data. We pick random row numbers with `sample()`, then keep those rows.

Because the choice is random, we first use `set.seed()`. This makes the “random” choice the same every time you run it, so your work stays **reproducible**.

```

set.seed(123)                # makes the random choice repeatable
sub_rows <- sample(1:nrow(survey), 10) # pick 10 random row numbers
subsample <- survey[sub_rows, ]      # keep those rows
nrow(subsample)
#> [1] 10

```

Your turn

Create your own subsample of **15** people in an object called `my_sample`. First write `set.seed(1)`, then use `sample()` and square brackets as above. Check it has 15 rows with `nrow()`.

7 Describing your sample

Now that the data is clean, we describe it. This is the first step of almost every study: who is in the sample, and how are the values spread out?

7.1 Counting categories with `table()`

For categories (text variables), we count how often each value appears with `table()`.

```

table(survey$program_name)
#>
#> Pedagogy
#>      10

```

7.2 Numbers in short: mean and sd

For numbers, the average (`mean()`), the median (`median()`), and the standard deviation (`sd()`) give a quick picture.

```


mean(survey$self_efficacy) # average self-efficacy
#> [1] 33.23333
sd(survey$self_efficacy)   # how spread out the scores are
#> [1] 6.043996

```

7.3 Histograms with hist()

A **histogram** shows how a numeric variable is distributed. It puts the values into bars: tall bars mean many people have those values.

```
hist(survey$self_efficacy,  
     main = "Self-efficacy scores", # title  
     xlab = "score")               # label for the x-axis
```




Your turn

Make a histogram of **age**. Give it a clear title and an x-axis label. Then, print the mean and median of **age**, are they similar or different? Argue why mean and median are similar or different, and relate it to the histogram you just plotted.

7.4 Boxplots with boxplot()


A **boxplot** is another way to show the distribution of a numeric variable. The thick line is the middle value (the median), the box holds the middle half of the people (the interquartile range), and the lines reach out to the smallest and largest typical values.

```
boxplot(survey$study_hours,  
        main = "Study hours per week",  
        ylab = "hours")
```



Boxplots are useful when you **compare groups**. We do that with a tilde ~, which you read as “by”. The variable before the ~ is the number; the variable after it is the group.

```
boxplot(study_hours ~ program_name,  
        data = survey,  
        main = "Study hours by programme",  
        ylab = "hours per week")
```



Now you can see whether students in some programmes study more than in some others in your sample.

Your turn

Make a boxplot of `self_efficacy` by `program_name`. Which programme has the highest middle value? What can you conclude from it?

8 Saving your cleaned data

You changed the data, so save the clean version under a **new** name. This keeps the original file safe. We use `write.csv()` and add `row.names = FALSE` so R does not add an extra column of row numbers.

```
write.csv(survey, "data/class_survey_clean.csv", row.names = FALSE)
```

Your turn

Save your cleaned `survey` as `class_survey_clean.csv` in your `data` folder. Then check in the **Files** pane (bottom-right) that the new file is there. Also, save your script in `script` folder.

9 Recap: what you can now do

Last week you handled single numbers and vectors. You can now work with a whole dataset. You can:

- load a CSV file with `read.csv()` and store it as a data frame;
- reach a single column with `$`, and a single cell with `[row, column]`;
- inspect data with `dim()`, `names()`, `head()`, `str()`, `summary()`, and `View()`;
- filter rows with a condition;
- make new variables and remove ones you do not need;
- recode values into clear labels and put numbers into categories;
- take a reproducible random subsample with `set.seed()` and `sample()`;
- describe your sample with `table()`, `mean()`, `sd()`, `hist()`, and `boxplot()`;
- save a cleaned dataset with `write.csv()`.

A small cheat sheet to keep:

You want to...	You write...
Load a CSV	<code>read.csv("data/file.csv")</code>
View a data frame	<code>View(df)</code>
Use one column	<code>df\$age</code>
Rows and columns count	<code>nrow(df)</code> , <code>ncol(df)</code>
First rows / structure	<code>head(df)</code> , <code>str(df)</code>
Summary of all variables	<code>summary(df)</code>
Keep some rows (filter)	<code>df[df\$age >= 18,]</code>
Make a new variable	<code>df\$new <- df\$x / 7</code>
Remove a variable	<code>df\$old <- NULL</code>
Recode a value	<code>df\$x[df\$x == 1] <- "label"</code>
Category from a rule	<code>ifelse(df\$x >= 32, "higher", "lower")</code>
Random subsample	<code>df[sample(1:nrow(df), 10),]</code>
Count a category	<code>table(df\$group)</code>
Histogram	<code>hist(df\$x)</code>
Boxplot by group	<code>boxplot(y ~ group, data = df)</code>
Save a CSV	<code>write.csv(df, "data/out.csv", row.names = FALSE)</code>

Well done. Keep your materials in the `research_methods` folder, we build directly on them next week.
